



# Smart Contract Vulnerability Detection Using Machine Learning

A Pattern-Based Approach to Label Generation and Classification

Lucas Kim

Course: CMSI 5350

# Why Smart Contract Security Matters

- Smart contracts manage **billions of dollars** in decentralized applications.
- Code-level vulnerabilities are the **primary cause** of blockchain exploits.
- Manual security auditing is expensive (\$10k–\$100k per contract) and time-consuming.
- Critical need for **automated**, **scalable** vulnerability detection systems.



# Research Question & Objectives



## Research Question

Can machine learning models effectively detect vulnerabilities in smart contracts using features extracted directly from source code?



## Primary Objectives

- F1-score  $\geq 0.85$
- False Positive Rate  $< 0.10$
- Extract comprehensive features
- Evaluate multiple ML models

# The Labeling Problem: The Core Blocker

## The Core Issue

Dataset contains 42,908 real contracts but **no vulnerability labels**.

Without labels, supervised learning is impossible.

## Why It Matters

Existing tools fail: Slither fails on 99% of contracts due to compilation errors. Manual labeling would take weeks.

⚠ Synthetic labels yield  $F1 < 0.30$



## Unlabeled Data

42,000+ Files  
0 Labels

# Dataset Description & Split

## Ethereum Smart Contract Dataset

- 📄 **Total Available:** 42,908 contracts
- 📅 **Processed:** 2,000 contracts (subset)
- 📅 **Time Period:** 2020-2025

## Stratified Split (80/20)

Random State 42, 'Balanced' Weighting



# The Breakthrough: Hybrid Detection

Transforming an unsupervised problem into a supervised one by generating labels based on code patterns.



## Tier 1: Slither Analysis

Attempts compilation and static analysis.

**~1% Success**

Most contracts fail to compile.



## Tier 2: Pattern-Based

Works directly on source code without compilation.

**100% Success**

Ensures every contract is labeled.

# 9 Major Vulnerability Patterns

## High Severity (+2 Score)

- **Reentrancy:** External calls followed by state updates.
- **Unchecked External Calls:** Calls without validation.
- **Delegatecall Usage:** Execution in calling context.

## Medium Severity (+1 Score)

- tx.origin usage
- Selfdestruct
- Integer overflow
- Uninitialized storage
- Low-level calls
- No access control

```
// Reentrancy Example

function withdraw() public {
    // 1. External Call
    msg.sender.call.value(amount)("");


    // 2. State Update (Too Late)
    balances[msg.sender] = 0;
}
```

⚠ Danger: Attacker can re-enter before balance is set to 0.

# Label Generation Implementation



# 71 Comprehensive Features



## 1. Opcode Frequency (35)

CALL, DELEGATECALL, SSTORE, SLOAD...

## 2. AST Node Features (21)

FunctionDefinition, ModifierDefinition, Control Structures...

## 3. Control Flow (5)

Nesting depth, Loop counts, Complexity...

## 4. Code Metrics (7)

LOC, Character count, Function count...

# | Model Training Configuration

## Logistic Regression

Baseline with L1/L2 regularization.



## Random Forest

Ensemble of decision trees.  
Tuned for depth and split.



## XGBoost

Gradient boosting. Tuned for learning rate and subsample.

5-Fold CV GridSearchCV Class Weight: Balanced

# Model Performance (Real Labels)

Model	F1-Score	Precision	Recall	FPR
Logistic Regression	0.9607	0.9853	0.9373	0.0354
<b>Random Forest (Best)</b>	<b>0.9825</b>	0.9894	0.9756	<b>0.0265</b>
XGBoost	0.9789	0.9859	0.9721	0.0354

**Target Exceeded**

F1 Score: 0.98 vs Target 0.85

**3x Better FPR**

0.0265 vs Target 0.10

# Best Model: Random Forest

## Key Strengths


- **Highest F1-score:** 0.9825
- **Lowest FPR:** 0.0265
- Balanced Precision (0.989) and Recall (0.976).
- Robust and ideal for production deployment.

Confusion Matrix (Test Set: 400)

	Pred: 0	Pred: 1
Act: 0	110 TN	3 FP
Act: 1	7 FN	280 TP

# Impact of Real Labels

Label quality is more critical than model sophistication.



# Technical Challenges Overcome



## Missing Labels

Solved via Pattern-based detection system.



## Compilation Errors

Solved via Source-code analysis (100% coverage).



## Class Imbalance

Solved via Stratified splits & balanced weighting.



## ID Matching

Standardized formats (99.4% match rate).

# Project Achievements



## Targets Exceeded

- ✓ F1-score: **0.96-0.98** (Target:  $\geq 0.85$ )
- ✓ FPR: **0.03** (Target:  $< 0.10$ )

## Contributions

- Pattern-based label generation system
- Hybrid Slither + pattern matching approach
- End-to-end production pipeline



## Key Findings & Impact

"We transformed an **unsupervised problem**  
into a **supervised one.**"

By solving the labeling problem through pattern-based detection,  
we enabled standard ML models to achieve state-of-the-art performance, proving that  
**data quality > model complexity.**

# Limitations & Future Work



## Model Improvements

Deep Learning (LSTM, Transformers), Ensemble methods.



## Feature Engineering

Graph-based features (CFG, DFG), Semantic embeddings.



## Production

Real-time API, Model versioning, SHAP explainability.

# Summary

The project's success hinged on solving the labeling problem.

Pattern-based detection enabled us to achieve:

F1 Score

**0.98**

FPR

**0.02**

Improvement

**+989%**



Thank you for your attention.

Lucas Kim | CMSI 5350